

LADRA: Log-based Abnormal Task Detection and Root-cause Analysis in Big Data Processing with Spark

Siyang Lu^a, Xiang Wei^{a,d}, Bingbing Rao^a, Byungchul Tak^b, Long Wang^c,
Liqiang Wang^a

^a*Dept. of Computer Science, University of Central Florida, Orlando, FL, USA*

^b*Dept. of Computer Science and Engineering, Kyungpook National University, Korea*

^c*IBM T.J. Watson Research Center, Yorktown Heights, NY, USA*

^d*School of Software Engineering, Beijing Jiaotong University, China*

Abstract

As Cyber Physical Systems (CPS) are being widely adopted by many domains, massive data generated by CPS become more reliant on the parallel computing platforms for analysis, wherein Spark is one of the most widely used big data frameworks. Spark's abnormal tasks may cause significant CPS performance degradation, and it is extremely challenging to detect and diagnose the root causes. To that end, we propose an innovative tool, named LADRA, for log-based abnormal tasks detection and root-cause analysis using Spark logs. In LADRA, a log parser first processes raw log files to convert into structured format and extract features. Then, a detection method is proposed to detect where and when abnormal tasks happen. After that, we extract factors based on these features in order to analyze root causes. Finally, we leverage General Regression Neural Network (GRNN) to identify root causes for abnormal tasks. The likelihood of reported root causes are presented to users according to the weighted factors by GRNN. LADRA is an off-line tool that can accurately analyze abnormality without extra monitoring overhead. Four potential root causes, *i.e.*, CPU, memory, network, and disk I/O, are considered. We have tested LADRA atop of three Spark benchmarks by injecting aforementioned root causes. Experimental results show that our proposed approach is more

Email addresses: siyang@knights.ucf.edu (Siyang Lu), lwang@cs.ucf.edu (Liqiang Wang)

accurate in the root cause analysis than other existing methods.

Keywords: Spark, Log Analysis, Abnormal Task, Root Cause

1. Introduction

Rapid growth of massive data size brings great challenges in Cyber Physical Systems (CPS), which integrate many analytic infrastructures as their components to analyze data autonomously. Several parallel computing frameworks that follows the MapReduce [1] paradigm are widely-used in real-world CPS applications to handle batch and streaming data, such as Hadoop [2] and Spark[3]. Among these, Spark has recently gained wide-adoption in CPS. Different from the Hadoop framework, Spark supports a more general programming model, in which an in-memory technique, called Resilient Distributed Dataset (RDD) [4], is used to store the input and intermediate data generated during computation stages.

While Spark is highly successful in CPS domain for data analytics, it could suffer from significant performance degradation under the existence of abnormal tasks. A task is considered *abnormal* if it shows significant delay in comparison with other tasks within the same stage. A few causes of such performance degradation can be due to ineffective coding, resource contention, and data locality problems [5, 6, 7, 8].

To mitigate such performance problems, Spark employs a speculation [9] mechanism to detect stragglers during runtime, in which slow tasks are re-scheduled after marked as stragglers. Spark performs speculative execution of tasks until a specified fraction (defined by `spark.speculation.quantile`, which is 75% by default) of tasks is completed. Spark identifies stragglers by checking whether the running tasks are much slower (*e.g.*, 1.5 times, by default) than the median of all successfully completed tasks in the current stage. However, speculation mechanism cannot detect all stragglers and does not provide the root causes of degraded performance. In addition, monitoring tools are usually heavy-weight and cause significant overhead, which may impact the

performance of Spark even for normal executions. Therefore, abnormal task detection and root cause analysis still remain grand challenges.

30 This paper proposes LADRA, an off-line tool for log-based abnormal tasks detection and root-cause analysis for big data processing with Spark. LADRA detects abnormal tasks by examining features extracted from logs and analyzes them to find root causes via a neural network model. Specifically, our proposed approach adopts a statistical spatial-temporal analysis for Spark logs, which
35 consists of Spark execution logs and JVM garbage collection (GC) logs related to resource usage. LADRA’s abnormal task detection method is more effective than Spark speculation, as all Spark stages are considered and abnormal tasks happened in any life span could be detected. Moreover, Spark’s report could be inaccurate because Spark uses only fixed amount of finished task duration to
40 speculate the unfinished tasks. Our approach reports the likelihood of each potential root cause, which can be leveraged by users to tune resource allocations and reduce the impact of abnormal tasks. For instance, in one of our experiments, LADRA reports that abnormal tasks are caused by 80% network and 20% CPU issues on victim nodes, users may check the network condition first,
45 then tune CPU usage accordingly. There are four major root causes for task abnormalities: CPU, memory, network, and disk I/O, all of which are considered by this paper.

We make the following contributions in this work.

- An abnormality detection method is proposed that can accurately locate
50 where and when abnormal task executions happen by analyzing the Spark logs.
- We report on the 22 log features and 7 factors that are effective in exposing the degree of abnormality from the analysis of Spark log and GC log.
- A neural network-based analysis method is proposed, which is more ac-
55 curate and provides the ranked likelihood for true root causes to better understand the performance problems at hand and to tune the Spark settings.

The rest of the paper is organized as follows. Section 2 introduces the background knowledge of Spark and its applications in CPS, and surveys the related work. Section 3 gives an overview of our approach. Section 4 illustrates the feature extraction from Spark logs and abnormal task detection based on these features. Section 5 presents factor synthesization for root cause analysis. Section 6 describes the details of root cause analysis using GRNN. Section 7 shows our experimental results by evaluating our approach on several widely used benchmarks. Section 8 summarizes our method and discusses its limitations and future work.

2. Related work and background

In this section, we give brief background of Spark scheduling mechanisms and its log structures. Then, we review related work in the area of the root cause analysis for big data platforms.

This paper significantly extends our previous approach [10], a statistical method for detecting task abnormalities and analyzing the root causes. In this paper, we propose a more effective and efficient approach that leverages the General Regression Neural Network (GRNN). Our previous approach diagnoses root causes by applying weights for each factor. One major problem with such a rule-based weight calculation approach was that it may cause false positives. In this paper, a GRNN-based method is applied in order to avoid the ad-hoc factor selection and weight computing.

2.1. Spark architecture and its log structure

Spark architecture: Apache Spark is an in-memory parallel computing framework for large-scale data processing. Moreover, to achieve the scalability and fault tolerance, Spark introduces resilient distributed data set (RDD) [4], which represents a read-only collection of objects partitioned across a set of machines that can be rebuilt if a partition is lost. As shown in Figure 1, Spark cluster consists of one master node and several slave nodes, named as workers,

which may contain one or more executors. When a Spark application is submitted, the master will request computing resource from the resource manager based on the requirement of the application. When the resource is ready, Spark scheduler distributes tasks to all executors to run in parallel. During this process, the master node will monitor the status of executors and collect results from worker nodes.

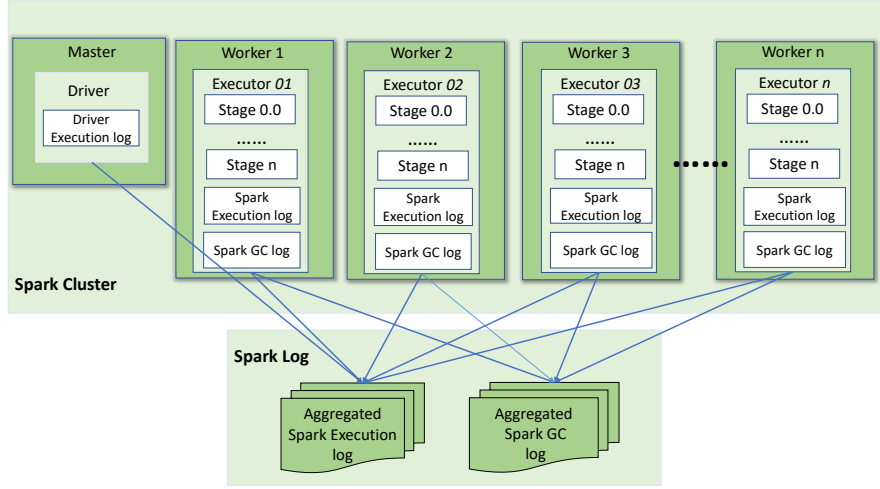


Figure 1: Spark workflow and log files

Spark logs include execution logs and JVM GC logs. Spark driver (master node) collects the information of all executors (*i.e.*, driver log), and each executor records the status of tasks, stages, and jobs within the executor (*i.e.*, execution log). Besides these logs, Spark JVM Garbage Collection (GC) logs are used by our analysis, which are the output from two output channels, `stderr` and `stdout`. When an application is finished, we use a parser to collect all Spark logs and aggregate them into two different categories: execution log and GC log automatically. An example is shown in Figure 2.

Spark uses “log4j”, a JAVA logging framework, as its logging framework. Spark users can customize “log4j” for their own logging by changing configuration parameters, such as log level, log pattern, and log direction. In this paper,

```

17/02/22 21:04:02.259 INFO
TaskSetManager: Starting task 12.0 in stage 1.0 (TID 58, 10.190.128.101, partition 12,
ANY, 5900 bytes)
17/02/22 21:04:02.259 INFO
CoarseGrainedSchedulerBackend$DriverEndpoint: Launching task 58 on executor id: 1
hostname: 10.190.128.101.
17/02/22 21:04:02.276 INFO
TaskSetManager: Finished task 1.0 in stage 1.0 (TID 47) in 14075 ms on 10.190.128.101
(1/384)

```

Figure 2: Spark execution log example.

we use the default configurations in “log4j”. As shown in Figure 2, each line of Spark execution log contains four types of information: *timestamp* with ISO
105 *format*, *logging level* (*e.g.*, INFO, WARNING, or ERROR), *related class* (which class prints out this message) and message content. A message content contains two main kinds of information, constant keywords (*e.g.*, **Finished task in stage TID in ms on**), and variables (*e.g.*, 1.0 1.0 47 14075..).

```

GC (Allocation Failure)
PSYoungGen: 95744K->9080K(111616K)
95744K->9088K(367104K), 0.0087250 secs] [Times: user=0.03, sys=0.01, real=0.01 secs]

```

Figure 3: Spark GC log example.

During the execution of a spark application, JVM monitors memory usage
110 and outputs the status information to GC logs when garbage collection is invoked. GC logs report two kinds of memory usage: heap space and young generation space, where young generation space is a part of heap memory space to store new objects. Figure 3 shows an example of Spark JVM GC log, where
115 “Allocation Failure” invokes this GC operation, and “PSYoungGen” shows the usage of young generation memory space. In “95744K->9080K(111616K)”, the first numeric is the young space before this GC happens, and the second one is the young space after this GC, the last one is the total young memory space. Similarly, “95744K->9088K(367104K)) illustrates heap memory instead of young generation space.

120 *2.2. Related work*

CPS applications have been widely used in people’s daily activities, such as smart home, smart factory, and smart city. Furthermore, being combined with the cloud computing framework, CPS is rapidly growing in size and becoming smarter [11, 12]. Cloud computing and big data platforms turn out to be
125 key components in CPS. As these platforms grow bigger, abnormality becomes more common. Thus, root cause analysis of abnormal performance becomes a necessity to maintain the reliability of CPS.

There are several categories of the root causes for the abnormal performances. Ananthanarayanan *et al.* [13] identify three categories of root causes
130 for Map-Reduce outliers: the key role cause is machine characteristics (resource problems), the other two causes are network and data skew problem. Ibidunmoye *et al.* [14] depict that four root causes may cause bottlenecks, which are system resource, workload size, platform problems, and application (buggy codes). Garbageman *et al.* [15] analyze around 20-day cloud center data and
135 summarize that most common root cause in cloud center of abnormal occurrence is server resource utilization, and data skew problems only take 3% of total root causes. According to the above studies on real world experiment, the primary root causes of abnormal tasks are machine resources, which includes CPU, memory, network, and disk I/O. Therefore, in our paper, we consider
140 the only the four main root causes, and ignore data skew and ineffective code problems.

Statistical and machine learning techniques are promising approaches in the root causes analysis, and their combination has been widely used in the parallel computing area to solve performance degradation problem caused by abnormal
145 executions. Abnormality detection and analysis using this approach can be categorized largely into online and offline approaches.

The online detection strategy is invoked during the executions of applications. For example, both Spark and Hadoop provide online “speculation” [9], which is a built-in component for detecting stragglers statistically. Although
150 it can detect stragglers during runtime, it does not offer the root causes. In

addition, the speculation is often inaccurate, *i.e.*, it may raise too many false alarms [16]. Chen *et al.* [17] propose a tool called Pinpoint that monitors the execution and uses log traces to identify the fault modules in J2EE applications via standard data mining approaches. A stream-based mining algorithm for
155 online anomalies prediction is presented by Gu *et al.* [18]. Ananthanarayanan *et al.* [13] design a task monitoring tool called Manrti, which can cut outliers and restart tasks in real time according to its monitoring strategy.

The offline strategy can be combined with the online one. Garraghan *et al.* [15] propose an empirical approach to extract execution paths for straggler
160 detection by leveraging an integrated off-line and online model. Some machine learning approaches are also leveraged in predicting system faults using logs and monitoring data, which is very similar to the root cause analysis. Fulp *et al.* [19] leverage a sliding window to parse system logs and predict failures using SVM. Yadwadkar *et al.* [20] propose an offline approach that works with resource
165 usage data collected from the monitoring tool Ganglia [21]. It leverages Hidden Markov Models (HMM), which is a liner machine learning approach.

Nevertheless, monitoring data may not be always accessible from the user side, due to the fact that the monitoring tools are hard to install and tune. Hence, some studies focus on the off-line strategy by analyzing logs instead of
170 monitoring [22, 23]. For example, Tan *et al.* [24] introduce a pure off-line state machine tool called SALSA, which simulates data flows and control flows in big data systems, and leverages Hadoop’s historical execution logs. Chen *et al.* [25] propose a self-adaptive tool called SAMR, which adds weights for calculating each task duration according to historical data analysis. Xu *et al.* [26] use an
175 automatic log parser to parse source code and combine PCA to detect anomaly, which is based on the abstract syntax tree (AST) to analyze source code and uses machine learning to train data.

There are some off-line approaches that analyze both log files and monitoring data to identify abnormal events. Aguilera *et al.* [27] propose two statistical
180 methods to discover causal paths in distributed system by analyzing historical log and monitoring data from the traces of applications. The most closely re-

lated work to our approach is BigRoots [28], which detects stragglers by Spark speculation and analyzes the root causes by extracted features. It leverages experience rule to extract features for each task from application log and monitoring data. However, the threshold in Spark speculation is not proper to detect abnormal tasks. In addition, BigRoots considers only the features for each individual task, which can not capture the status change of the cluster, thus such a rule-based method is very limited. In our method, we choose the combination of features to create the factors presenting the status change of the whole cluster, and a GRNN technique is leveraged instead of a rule-based statistical approach to avoid the limits.

3. Overview of LADRA’s approach

Although Spark logs are informative, there is rarely a direct information about the root cause of the abnormal tasks. This renders simple keyword-based log search ineffective for diagnosing the abnormal tasks. It motivates us to design LADRA to help users detect abnormal tasks and analyze their root causes. An overview of our approach is depicted in Figure 4, which contains five primary components: log preprocessing, feature extraction, abnormal task detection, factor extraction, and root cause analysis.

1. *Log preprocessing*: Spark log contains a large amount of information. In order to extract useful information for analysis, we first collect all Spark logs, including execution logs and JVM GC logs, from the driver node and all worker nodes. Then, we use a parser to eliminate noisy and trivial logs, and convert them into a structured data.
2. *Feature extraction*: Based on the Spark scheduling and potential abnormal task occurring conditions, we quantify the data locality feature with a binary number format. Then, we screen structured logs and select appropriate data in three kinds of feature sets: execution-related, memory-related, and system-related data. Finally, we store them into two numerical matrices: execution log matrix and GC matrix.

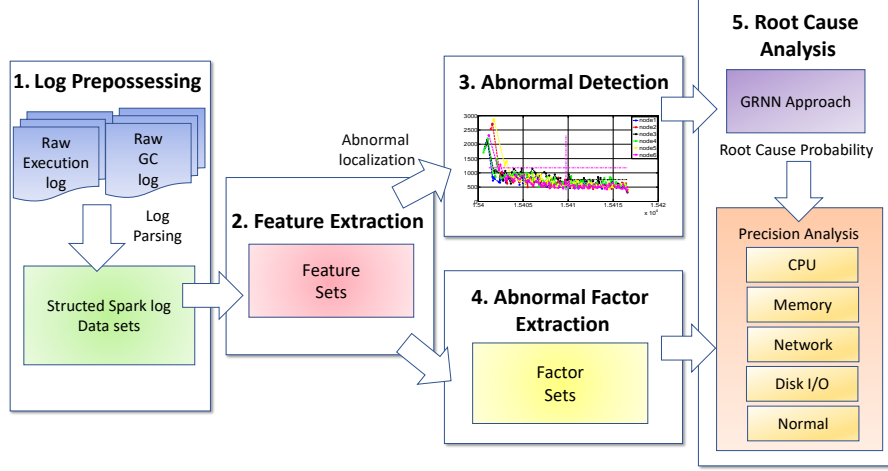


Figure 4: Workflow of LADRA

3. *Abnormal detection*: We implement a statistical abnormal detection algorithm to detect where and when the abnormal tasks happen based on the analysis of execution-related feature sets. This detection method determines the threshold by calculating the standard deviation of task duration and use it to detect abnormal tasks in each stage from Spark logs, which is introduced in details in Section 4.
4. *Abnormal factor extraction*: According to our empirical case study, we combine special features to synthesize two kinds of factors, the speed factor and the degree factor, which can describe the status of each node in the whole cluster. Those factors are used by our root cause analysis method. The factors are introduced in details in Section 5.
5. *Root cause analysis*: We propose a General Regression Neural Network (GRNN) based approach for our root-cause analysis, in which probability result can be calculated more accurately than our previous statistical work. Our experiments show that the GRNN-based approach has more accurate results than existing approaches, which are introduced in details in Section 6.

Table 1: Extracted features for abnormal task detection

Feature Category	Feature Name		
Execution related	Task ID	Job ID	Task duration
	Stage ID	Job duration	Data locality
	Host ID	Stage duration	Timestamp
	Executor ID	Application execution time	
Memory related	GC time	After young GC	After Heap GC
	Full GC time	Before young GC	Before Heap GC
	Heap space	GC category	
System related	Real time	CPU time	User time

4. Log feature extraction and abnormal task detection

4.1. Log feature extraction

230 When an abnormal task happen, it usually does not cast any warnings or error messages. As Spark does not directly reveal any information about abnormal tasks, it is a very challenging problem to detect these problems. Our approach starts from understanding Spark scheduling strategy, then extracts features associated with CPU, memory, network, disk I/O to build a feature matrix, which
 235 reflects the whole cluster’s status. These features can be classified into three categories: execution-related, memory-related, and system-related, as shown in Table 1.

The execution-related features are extracted from Spark execution logs, including (1) the ID number of each task, stage, executor, job, and host, (2)
 240 the duration of each task, stage, and job, (3) the whole application execution time, (4) the timestamp for each event, and (5) data locality. Spark GC logs represent JVM memory usage of each executor in workers, from which we can extract memory-related features such as heap usage, young space usage before GC, young space usage after GC. In addition, system related features can be
 245 also extracted from GC logs, such as real time, system time, and user time.

4.2. Abnormal task detection

Our abnormal task detection is based on the extracted feature sets. In order to eliminate the false negative problem in the Spark speculation’s detection mentioned in Section 1, a more robust approach is designed to locate where and when an abnormal case happens, which includes the following two steps.

Step-1: Comparing task duration on inter-node: One basic rule for abnormal task identification is that the duration of abnormal task is relatively much longer than the duration of normal tasks (long tail). In the existing approaches for abnormal detection, both Hadoop and Spark use speculation, and [15] uses “mean” and “median” to decide the threshold. However, to seek a more reasonable abnormal detection strategy, we consider not only the mean and median of task running time, but also the distribution of the whole tasks’ duration, namely the standard deviation. In this way, we can get a macro-awareness on the task duration, and then based on the distribution of data, a more reasonable threshold can be determined to differentiate abnormal from the normal ones.

We compare the duration of tasks in the same stage but across different nodes (inter-node). Let $T_task_{i,j,k}$ denote the execution duration of task k in stage i on node j . And let avg_stage_i denote the average execution time of all tasks, which run on different nodes in the same stage i .

$$avg_stage_i = \frac{1}{\sum_{j=1}^J K_j} \left(\sum_{j=1}^J \sum_{k=1}^{K_j} T_task_{i,j,k} \right) \quad (1)$$

where J and K_j are the total node numbers and total task numbers in node j , respectively.

To determine a more proper threshold, we leverage the standard deviation of tasks duration in stage j of all nodes, which is denoted by std_stage_i , and λ is a threshold parameter used in Spark speculation, which is 1.5 by default. Thus, abnormal tasks can be determined by the following conditions:

$$task_k \begin{cases} abnormal & T_task_k > avg_stage_i + \lambda * std_stage_i \\ normal & otherwise. \end{cases} \quad (2)$$

Step-2: Locating abnormal task happening: After the first step, all tasks are classified into “normal” and “abnormal”, the time line are labeled as a vector with binary number (*i.e.*, 0 or 1, which denotes normal and abnormal, respectively). To smooth the outliers (*e.g.*, 1 appears after many continuous 0) inside each vector, which could be an abrupt change but inconsistent abnormal case, we empirically set a sliding window with the size of 5 to scan this vector. If the sum of numbers inside the window is larger than 2, the number in the center of the window will be set to 1, otherwise 0.

The next step is to locate the start timestamp and end timestamp of the current abnormal task. Note that, as Spark logs record the task finishing time but not the start time, so we locate the real abnormal task’s start time as the recorded task finishing time minus its duration. Moreover, to detect abnormal tasks in each stage, we classify tasks into two sets. One is for the initial tasks whose start timestamp are the beginning of each stage, as these tasks often have more overhead (such as loading code and Java jar packages), and they usually last much longer than the following tasks. Another set consists of the rest tasks. Our experiments show that this classification inside each stage can lead to a much accurate abnormal threshold. In this way, our abnormal detection method can not only detect whether abnormal tasks happen, but also locate where and when they happen.

Figure 5 shows abnormal detection process in our experiment for Spark WordCount under CPU interference. Figure 5 (a) and (b) are two stages inside the whole application. Moreover, inside each of the stage, purple dot-line is the abnormal threshold determined by Eq. (1), and the black dot-line indicates the threshold calculated by Spark speculation. For all tasks within a certain stage, the duration longer than the threshold are determined as abnormal tasks; otherwise, they are normal. Figure 5 (d) uses memory-related feature to display memory occupation along the execution of its corresponding working stages.

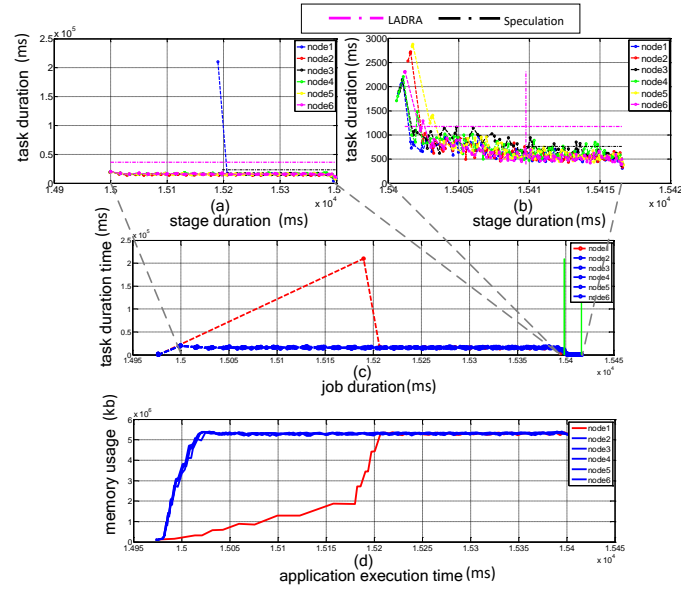


Figure 5: Abnormal detection under CPU interference in the experiment of WordCount: (a) Abnormal detection result in Stage-1. (b) Abnormal detection result in Stage-2. (c) Execution-related feature visualization for abnormal detection in the whole execution. (d) memory-related feature visualization for abnormal detection in the whole execution.

300 As we mentioned before, the data skew problem is not within our four considered root causes. Therefore, in the real analysis, those abnormal tasks caused by data skew should be eliminated as a noise. Data skew tasks can be easily detected by checking data locality features (*i.e.*, target data is not on the current node) combined with task duration features from execution logs.

305 5. Factor extraction for root-cause analysis

To look for the root causes of abnormal tasks, we introduce abnormal factors, which are meaningful synthesizations of features based on empirical study on the 22 features in Spark log matrix and GC matrix. Namely, those factors are normalized features that present status change of the whole cluster, not only for assessing individual components, such as task and stage, but also a series of abnormal tasks, which may be generated by continuous interference affecting the cluster. In normal cases, each factor should be close to 1; otherwise, it implies an abnormal case. In our factors' definition, j denotes the j th node, J presents a set of nodes; i indicates the index of stage, I is a set of stages; k denotes a task, K is a task set; n stands for a GC record, N is a GC record set. 310 All factors used to determine root causes are listed as below.

Degree of Abnormal Ratio (DAR) describes the degree of imbalanced scheduling of victim nodes, due to the fact that the victim nodes will be scheduled with fewer tasks than other normal nodes. For example, as shown in Figure 320 6, CPU interference can cause fewer tasks (red dots) scheduled at a victim node (node1) than normal nodes. Eq. (3) illustrates the degree of abnormal ratio in a certain stage. Therefore, the factor DAR indicates that the number of tasks in intra-node on a certain stage is considered for abnormal detection.

$$DAR = \frac{\frac{1}{J-1}((\sum_{j=1}^J k_j) - k_{j'})}{k_{j'}} \quad (3)$$

where k_j denotes the number of tasks on node j , and J is the total number of nodes in the cluster. Here, we assume that node j' is abnormal. 325




Figure 6: CPU interference injected after Sorting application has been submitted for 60s, and continuously impacts for 120s.

Degree of Abnormal Duration (DAD) is used to measure the average task duration, as the abnormal nodes often record longer task duration.

$$DAD = \frac{avg_node_{j'}}{\frac{1}{J-1}((\sum_{j=1}^J avg_node_j) - avg_node_{j'})} \quad (4)$$

where avg_node_j is defined as:

$$avg_node_j = \frac{1}{K_j}(\sum_{k=1}^{K_j} T_task_{i,j,k}) \quad (5)$$

Degree of CPU Occupation (DCO) describes the degree of CPU occupation by calculating the ratio between the wall-clock time and the real CPU time. In the normal multiple-core environment, “realTime” is often less than “sysTime+userTime”, because GC is usually invoked in a multi-threading way. However, if the “realTime” is bigger than “sysTime+userTime”, it may indicate that the system is quite busy due to CPU or disk I/O contention. We choose a max value across nodes as the final factor.

$$DCO = \max_{j \in J} (avg(\frac{realTime_{i,j}}{sysTime_{i,j} + userTime_{i,j}})) \quad (6)$$

Memory Change Speed (MCS) indicates the speed of memory usage change according to GC curve. Due to the fact that under CPU, memory, and disk I/O interference, the victim node's GC curve will vary slower than the normal nodes' GC curve, as shown in Figure 7. $start_a$ and $stable_a$ are the points of the start position (the corresponding memory usage at abnormal starting time) and the stable memory usage position, respectively. $start_b$ and $stable_b$ are the start and end positions of abnormal memory, respectively, which are obtained by analyzing logs introduced before. The motivation is that the interfered node gradually uses less memory than normal nodes under interference, as shown in Figure 7. Hence, we use the area under GC curve a in the whole cluster ($start_a$ of normal node) to calculate this factor, as shown in Eq. (7).

$$MCS = \frac{\int_{start}^{stable_a} f(x_a)dx_a}{\int_{start}^{stable_b} f(x_b)dx_b} \quad (7)$$




Figure 7: MCS in CPU interference injected after WordCount application has been submitted for 20s, and continuously impacts for 120s.

Abnormal Recovery Speed (ARS) measures the speed of abnormal task's recovery. Obviously, one Spark node often accesses data from other nodes, which leads to network interference propagation. It is both inter-node and intra-

node problem. In this way, we can detect network interference happening inside cluster, as shown in Figure 8, which is the location of our detected interference and shows that task duration will be affected by delayed data transmission. We leverage Eq. (8) to calculate this factor, where abn_prob_j indicates the ratio of abnormal that we detect for each node j inside that area. The reason that we use the product of abnormal ratio other than the sum of them is that only when all nodes are with a portion of abnormal, we identify them with a potential of network interference; if their sum is used, we cannot detect this joint probability. Meanwhile, the exponential is to make sure this factor is no less than 1. Hence, the phenomenon of error propagation will be detected and quantified by calculating this factor.

$$ARS = \exp(J * \prod_{j=1}^J abn_prob_j) \quad (8)$$




Figure 8: Network interference is injected after WordCount has been executed for 100s, and continuously impacts for 160s.

Degree of Memory Change (DMC) describes how much of memory usage changed during the execution in each node. In fact, when network bandwidth is limited, or the network speed slows down, the victim node gets affected

by that interference, and tasks will wait for their data transformation from other
 365 nodes. Hence, the tasks will pause or work very slowly, and data transfer rate
 becomes low, as shown in Figure 9. We leverage Eq. (9) to find the longest
 horizontal line that presents the conditions under which tasks' progress become
 tardy (*e.g.*, CPU is relatively idle and memory remains the same). In Eq. (9),
 $m_{j,n}$ indicates the gradient of memory changing in the n th task on node j .
 370 First, the max value of gradient is calculated for each GC point, denoted as
 m . Second, we make a trade-off between its gradient and the corresponding
 horizontal length to identify the longest horizontal line in each node. Then, to
 determine a relative value that presents the degree of abnormal out of normal,
 we finally compare the max and min among nodes with their max "horizontal
 375 factor" ($e^{-|m_{j,n}|} * (x_{j,n} - x_{j,n-1})$), where e is to ensure that the whole factor of
 b not less than 1.

$$DMC = \frac{\max_{j \in J} \{ \max_{n \in N} [e^{-|m_{j,n}|} * (x_{j,n} - x_{j,n-1})] \}}{\min_{j \in J} \{ \max_{n \in N} [e^{-|m_{j,n}|} * (x_{j,n} - x_{j,n-1})] \}} \quad (9)$$

where $m_{j,n} = \frac{y_{j,n} - y_{j,n-1}}{x_{j,n} - x_{j,n-1}}$.




Figure 9: Network interference is injected after WordCount has been executed for 30s, and continuously impacts for 160s.

Degree of Loading Delay (DLD) measures how much difference of loading duration on cluster nodes. Note that the initial task at the beginning of each stage always has a higher overhead to load data compared with the rest tasks. Similar to the factor DMC, instead of taking all tasks inside the detected stage into consideration, here, the first task of each node is used to replace the “*avg_node_j*”. Instead of taking all the tasks inside the detected stage into consideration, here, the first task of each node is used to replace the “*avg_node_j*” in Eq. (4). Formally, the equation is modified as Eq. (10) shows.

$$DLD = \frac{T_task_{i,j',1}}{\text{avg}_{j \in J}(T_task_{i,j,1})} \text{ where, } j' \notin J \quad (10)$$

6. Root cause analysis

6.1. Statistical rule based approach

We propose a statistical rule based approach for root cause analysis in [10]. As shown in Table 2, each root cause is determined by a combination of factors with specific weights.

The nodes with CPU interference often have a relatively lower computation capacity, which leads to less tasks allocated and longer execution time for tasks on it. Factors DAR and DAD are used to test whether the interference is CPU or not, because CPU interference can reduce the number of scheduled tasks and increase the abnormal tasks’ execution time. Factor DCO indicates the degree of CPU occupation, and CPU interference will slow down of the performance compared to normal cases. Factor MCS is used to measure memory changing rate, because CPU interference may lead memory change, thus the nodes become slow than other regular nodes.

For the network-related interference, because of its propagation, the nodes interfered earlier will often recover earlier, too. So our approach is to detect the first recovered node as the initial network-interfered node, and the degree ARS quantitatively describes the interference. When network interference occurs, tasks are usually waiting for data delivery (factor DMC).

Table 2: Related factors for each root cause

Factor	CPU	Mem	Network	Disk
DAR	✓	✓		
DAD	✓	✓	✓	
DCO	✓			✓
MCS	✓	✓		✓
ARS			✓	
DMC			✓	
DLD				✓

405 For the memory-related interference, when memory interference is injected into the cluster, we can even detect a relatively lower CPU usage than other normal nodes. Considering this, the task numbers (factor DAR) and task duration (factor DAD) are also added to determine such root causes with certain weights. Moreover, the memory interference will impact memory usage, and the factor MCS should be considered for this root cause detection.

410 To determine disk-related interference, we introduce the factor DLD to measure the degree of disk interference. The task set scheduled at the beginning of each stage could be affected by disk I/O. Therefore, these initial tasks on disk I/O interfered nodes behave differently from other nodes' initial tasks beginning tasks (factor DLD), CPU will become busy, and memory usage is different from other nodes'. Therefore, the memory changing rate (factor MCS) and CPU Occupation (factor DCO) are also used to determine such root causes.

420 After deciding the combination of factors for each root cause, we give them weights to determine root causes accurately as Eq. (11) shows. Here, all weights are between 0 and 1, and the sum of them for each root cause is 1. To decide the values of weights, we use classical liner regression on training sets that we obtained from experiments. Eq. (12) is proposed to calculate the final probability that the abnormal belongs to each of the root causes.

$$\begin{aligned}
CPU &= 0.3 * DAR + 0.3 * DAD + 0.2 * DCO + 0.2 * MCS \\
Memory &= 0.25 * DAR + 0.25 * DAD + 0.5 * MCS \\
Network &= 0.1 * DAD + 0.4 * ARS + 0.5 * DMC \\
Disk &= 0.2 * DCO + 0.2 * MCS + 0.6 * DLD
\end{aligned} \tag{11}$$

$$probability = 1 - \frac{1}{factor} \tag{12}$$

To sum up, the statistical rule based approach offers a reasonable result to
425 explain the its root causes probabilities. However it can not give a satisfied
result with higher precision for its classifying. Since the relationship between
factors is not simply linearly correlated, and we also changed old factor MCR
to a new factor MCS with AUC calculation instead of gradients calculation and
add it to our factor sets. From this point, a GRNN-based approach is proposed
430 for root cause analysis to consider no-linearly correlated relationship of new
factor set, and avoid human ad-hoc choosing and classification.

6.2. GRNN approach

In this paper, we propose a new neural network based model to automatically
calculate the probability of each root cause. We use a one-pass training neural
435 network, GRNN, to create a smooth transition and more accurate results.

In our previous work on root cause detection [10], which is briefed in Section
6.1, we use linear regression to identify the relationship among factors and root
causes. However, due to the complex relationship between hardware and soft-
ware that reflects the complexity of the relationship between input and output,
440 we believe that a non-linear model can do a better job. As we stated before,
the root cause detection work is better to be treated as a regression other than
a classification problem, after a bunch of attempts, we choose GRNN as a more
suitable choice for this work. In brief, the BP-class deep learning algorithms
may be vulnerable to the over-fitting problem especially when the dataset is
445 small [29], which is just the characteristic of our dataset. For some traditional

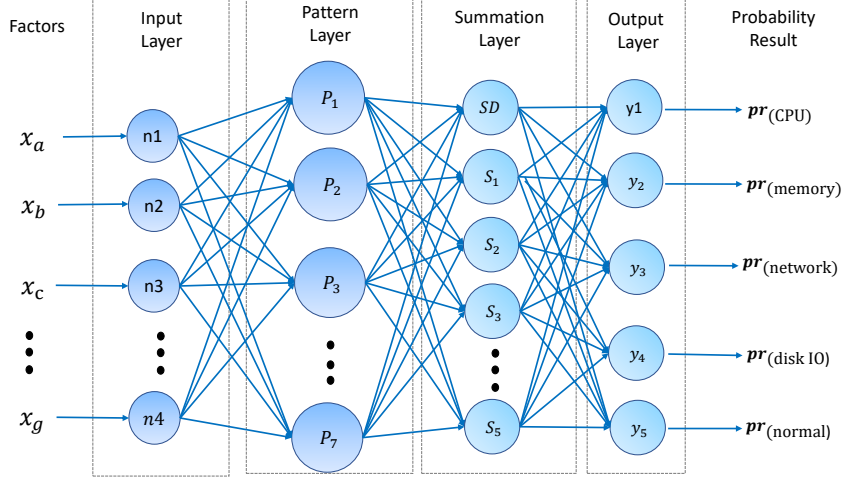


Figure 10: The architecture of our GRNN-based model for root-cause analysis.

data fitting algorithms, they suppose that the data obey some kind of distribution in advance, which can drastically affect the final result. As a non-parameter neural network model for data fitting, with its high efficiency and accuracy, GRNN is fully capable of dealing with our current problem. In addition, the experimental results demonstrate the effectiveness of GRNN compared with other attempts we have tried.

A representation of the GRNN architecture for our implementation of root cause identification is shown in Figure 10. Our model consists of four layers: input layer, hidden layer, summation layer, and output layer. For the input layer, it consists of 7 neurons, which indicates the dimension of input feature vector. After input layer, hidden layer is connected, which consists of neurons with the same size as input data, and the transfer function in hidden layer is defined in (14). In the following, summation layer is added, which contains two kinds of neurons: S-summation neuron (S) and D-summation neuron (SD), as defined in (13). SD neuron is used for calculating the arithmetic summation of pattern layer's output. The remaining S neurons weight summation for the

output of pattern layer. Transfer functions for two summation neurons are shown below. The i and j denote the number of input and number of output, respectively. The label (output layer) here is a 5-dimension one-hot vector with one indicating normal log and the rest four are injections. Due to probability
465 representation of root cause, after the output layer of GRNN, we add a softmax layer to convert the sum of 5-dimensional output to be 1.

$$Summations \begin{cases} SD = \sum_{i=1}^n (F_i), \text{ where } i = 1, 2, 3, \dots, n \\ S_j = \sum_{i=1}^n (y_{ij} F_i), \text{ where } j = 1, 2, 3, 4, 5 \end{cases} \quad (13)$$

where n is equal to 7, F_i is defined as:

$$F_i = e^{\frac{-(X-X_i)^T(X-X_i)}{2\sigma^2}} \quad (14)$$

$$y_j = \frac{S_j}{SD}, \text{ where } j = 1, 2, 3, 4, 5 \quad (15)$$

To sum up, GRNN can select a dominant weight for each of our factors, and
470 provide the root cause probability results with high accuracy.

7. Experiments

We evaluate LADRA on four widely used benchmarks and focus on the following two questions: (1) Can the abnormal tasks be detected? (2) What accuracy can LADRA's root cause analysis achieve? In the experiment, we
475 conduct a series of interference injections to simulate various scenarios that lead to abnormal tasks.

7.1. Setup

Clusters: We set up an Apache Spark standalone cluster with one master node (labeled by m1) and six slave nodes (labeled by n1,n2,n3,n4,n5,n6) based
480 on Amazon EC2 cloud resource. Each node is configured with type of "r3.xlarge" (24 virtual cores and 30GB of memory) and Ubuntu 16.04.9. We conduct a

bunch of experiments atop of Apache Spark 2.2.0 with JDK 1.8.0, Scala-2.11.11, and Hadoop-2.7.4 packages. Given that an AWS instance is configured with EBS by default, it is difficult for us to inject disk I/O interference. Hence, we set up
485 a 90G ephemeral disk for each instance and deploy a HDFS to store data.

Table 3: Benchmark resource intensity

	CPU	Memory	Disk I/O	Network
WordCount	✓		✓	✓
Sorting	✓		✓	✓
K-Means	✓			✓
PageRank	✓	✓		

Workload: In fact, some Spark applications may consume resources more intensively. According to previous studies on Spark performance [30], we choose four benchmarks built on Hibench [31] in our experiments: WordCount, Sorting, PageRank, K-means, which cover the domain of statistical batch application,
490 machine learning program, and iterative application. WordCount and Sorting are one-pass programs, K-means and PageRank are iterative programs. We characterize the benchmarks by resource intensive type and program type for underpinning our approach’s scalability. The resource intensity of each benchmark is shown in Table 3. The characteristics of four benchmarks are listed as
495 follows.

- WordCount is a one-pass program for counting how many times a word appears. We leverage RandomTextWriter in Hibench to generate 80G datasets as our workload and store it in HDFS. It is CPU-bound and disk-bound during map stage, then network-bound during reduce stage.
- 500 • Sorting is also a one-pass program that encounters heavy shuffle. The input data is generated by RandomTextWriter in Hibench. Sorting is disk-bound in sampling stage and CPU-bound in map stage, and its reduce stage is network-bound.
- K-means is an iterative clustering machine learning algorithm. We use
505 Hibench k-means to evaluate our approach. The workload is generated by

the k-means generator in SPARKBENCH, and is composed of 80 million points and 12 columns (dimensions). It is CPU-bound and network-bound during map stage.

- PageRank is an iterative ranking algorithm for graph computing. In order to analyze root causes of abnormal tasks with PageRank, we use Hibench PageRank as the testing workload, and generate eighty thousand vertices by Hibench’s generator as input datasets. It is CPU-bound in each iteration’s map stage, and network bound in each reduce stage.

7.2. LADRA interference framework

In order to induce abnormal tasks in the real execution for experiment, we design an interference framework that can inject four major resource (CPU, memory, disk I/O, and network) interference to mimic various abnormal scenarios. In order to simplify experiment, we apply all interference injection techniques only on node n1 for all test cases. In addition, for each injection, it will be launched during a time interval of 10 seconds and 60 seconds after the first spark job is initiated, and continue for 120 seconds to 300 seconds. Finally, when a test case is over, we recover all involved computing nodes to normal state by terminating all interference injections. Specifically, the following interference injections are used in our experiments:

- *CPU interference*: CPU Hog is simulated via spawning a bunch of processes at the same time to compete with Apache Spark processes. This injection causes CPU resource contention in consequence of limited CPU resource.
- *Memory interference*: Memory resource scarcity is simulated via running a program that requests a significant amount of memory in a certain time to compete with Apache Spark jobs, then we hold on this certain of memory space for a while. Thus, Garbage Collection will be frequently invoked to reclaim free space.

Table 4: LADRA’s abnormal task detection compares with Spark speculation’s approach in four intensive benchmarks, where TPR = True Positive Rate, FPR = False Positive Rate.

Abnormal tasks detection	LADRA		Spark speculation	
	TPR	FPR	TPR	FPR
WordCount	0.96	0.06	0.94	0.8
Sorting	0.96	0.16	0.96	0.7
K-Means	0.7	0.1	0.2	0.7
PageRank	0.6	0.517	0.9	0.48

- *Disk interference*: Disk Hog (contention) is simulated via leveraging “dd” command to continuously read data and write them back to the ephemeral disk to compete with Apache Spark jobs. It impacts both write and read speed. After the interference is done, we clear the generated files and system cache space.
- *Network interference*: Network scenario is simulated when network latency has a great impact on Spark. Specifically, we use “tc” command to limit bandwidth between two computing nodes with specific duration. In this way, the data transmission rate will be slowed down for a while.

7.3. Abnormal task detection

To evaluate LARDA, we compare LADRA’s detection with the Spark speculation. Each benchmark is executed 50 times without any interference injection, and 50 times under the circumstances of abnormal tasks. After that, we calculate the True Positive Rate (TPR) and False Positive Rate (FPR) results by counting the correct rate of each job classification as shown in Eq. (16) and Eq. (17). The comparison result is shown in Table 4.

As a build-in straggler detector, Spark speculation brings False Positive (FP) and True Negative (TN) problems in abnormal task detection. We compare LADRA with Spark speculation in details. For instance, Figure 11 shows one stage in a normal K-means execution, x-axis and y-axis present stage duration and task duration, respectively, and no abnormal tasks are detected by LADRA

Table 5: Root cause analysis result of LADRA’s GRNN approach, TPR = True Positive Rate,
P = Precision

GRNN	WordCount		Sorting		K-Means		Pagerank	
	TPR	P	TPR	P	TPR	P	TPR	P
CPU	1.0000	1.0000	1.0000	0.9400	0.8571	0.8350	0.9510	0.8261
Disk I/O	0.4500	0.4200	0.6790	0.8947	0.4233	0.6923	0.5400	0.8470
Network	1.0000	0.9545	1.0000	0.8529	0.6788	0.7300	0.6875	0.5640
Normal	0.9188	0.8378	0.9652	0.9243	0.7333	0.6857	0.5322	0.6400

(purple higher horizontal dash dotted line). However, Spark speculation (black lower horizontal dash dotted line) detects stragglers (area above the speculation line and beside red dotted vertical line) after 75% tasks (red dotted vertical line) finish. In this way, Spark speculation may delay the normal execution, as it will reschedule the stragglers to other executors. Moreover, Spark speculation will cause true negative problems as shown in Figure 6, because it only checks the 25% slowest tasks. As shown in Table 4, LADRA has a better accuracy in abnormal task detection than Spark speculation for all benchmarks. However, LADRA has lower accuracy on K-Means and PageRank than WordCount and Sorting. We find that under normal execution, most tasks in the map stage or sampling stage of K-Means and PageRank have an unexpected longer duration, because these benchmarks have many iteration stages, and tasks in those stages have data skew and cross-rack traffic fetching problems. LADRA cannot detect data skew problem within normal detection results. Too many such kinds of tasks with unexpected duration will cause LADRA to report false positives.

$$TPR = \frac{TP}{(TP + FN)} \quad (16)$$

$$FPR = \frac{FP}{(TP + TN)} \quad (17)$$




Figure 11: Abnormal task detection for K-means without interference injection.

7.4. LADRA's root cause analysis result

To test the accuracy of LADRA's GRNN approach for root cause analysis, we use cross validation strategy with 1/3 for test data and 2/3 for train data each time. Data in normal cases is also used in our training for improving the accuracy. In order to demonstrate the effectiveness of our approach, we run the GRNN 100 times and get the final accuracy result. We calculate the Precision (P) and True Positive Rate (TPR) for each detected root cause type by Eq. (18) and Eq. (16).

$$P = \frac{TP}{(TP + FP)} \quad (18)$$

We abandon memory root cause analysis in our experiments for three reasons. First, injecting significant memory interference into one node may cause the whole application to crash, as executors of Spark will fail if without enough memory. For instance, injected memory interference in PageRank benchmark not only causes Out-of-Memory (OOM) failures, but also makes executor keep quitting (executors are continuously restarted and fail). Secondly, memory interference does not work for non memory-intensive benchmarks. For instance, WordCount is not a memory-intensive program, and it will not evoke abnormal

tasks, even injecting significant memory interference. Thirdly, memory interference could also consume CPU resources, and may mislead GRNN’s classifying.

590 Table 5 summarizes the total P and TPR results of LADRA’s root cause analysis for four benchmarks. There are two issues to be noted. (1) LADRA has the highest CPU analysis precision (1.000 in CPU root cause analysis for WordCount) and higher network analysis precision (0.9545 in network root cause analysis for WordCount) results than disk I/O (0.4200 in disk I/O root cause
595 analysis for WordCount) for three reasons. First, all four benchmarks are CPU-intensive, and require large CPU resource for computing (map and sampling stages), and network resource to transfer data (reduce stages). Secondly, abnormal tasks have longer duration after CPU interference is injected, and the impact of network injection is significant (CPU stays idle). Thus, the synthesized factors demonstrate their effectiveness. Thirdly, as disk hog is injected by
600 leveraging a bunch of processes to read and write disk, it consumes not only disk I/O but also a certain of CPU resources. Therefore, disk I/O injections may be wrongly classified into other root causes (*e.g.*, CPU, network, or normal). (2) As shown by Table 5, LADRA is more precise on one-pass benchmarks than
605 iterative benchmarks, such as K-means and PageRank. The TPR of k-means and PageRank’s disk I/O is lower than the other two benchmarks. It is because that PageRank and k-means are not disk I/O-intensive benchmarks, if the intermediate data is small enough to be caught in memory, it will not use disk space. Therefore, the disk interference does not impact too much for these
610 benchmarks that have small size intermediate data. Moreover, wrong classification of other root causes in k-means and PageRank also impacts LADRA’s normal root cause classification, it causes more FP problems, or less TP. So the normal cases in k-means and PageRank also have lower precision and TPR.

To sum up, LADRA can analyze root causes via Spark log with high precision
615 and TPR for one-pass applications. However, there may be a few of limitations for LADRA to analyze root causes by only using Spark logs. Although Spark logs contain full information, but not so rich as monitoring data. It might be not possible to analyze all kinds of root causes by only leveraging log files. Some

root causes such as code failures, resource usages, and network failures, may rely
620 on monitoring tools. LADRA’s goal is to mine useful information and leverage
limited log information to analyze resource root causes without extra overhead.

8. Conclusions and future work

This paper presents LADRA, an off-line log-based root cause analysis tool to
effectively detect abnormal tasks of Spark. LADRA can identify abnormal tasks
625 by analyzing extracted features from Spark logs, which is more accurate than
Spark’s speculated straggler detection. In addition, LADRA analyzes the root
causes precisely using a GRNN-based method without additional monitoring.
The experimental results demonstrate that the proposed approach can accu-
rately locate abnormal and report their root causes for real-world benchmarks.

630 For the future work, we will consider more complex scenarios, such as mul-
tiple interference happen in parallel, to make our framework more robust for
root cause analysis. Moreover, we plan to design an automatic parser to ana-
lyze the source code of Spark applications, and obtain features to automatically
determine abnormal tasks and analyze root causes.

635 References

- [1] J. Dean, S. Ghemawat, Mapreduce: simplified data processing on large
clusters, *Communications of the ACM* 51 (1) (2008) 107–113.
- [2] Apache Hadoop website, <http://hadoop.apache.org/>.
- [3] Apache Spark website, <http://Spark.apache.org/>.
- 640 [4] M. Zaharia, M. Chowdhury, T. Das, A. Dave, J. Ma, M. McCauley, M. J.
Franklin, S. Shenker, I. Stoica, Resilient distributed datasets: A fault-
tolerant abstraction for in-memory cluster computing, in: *NSDI, USENIX*
Association, 2012.

- [5] H. Zhang, H. Huang, L. Wang, Mrapid: An efficient short job optimizer on hadoop, in: Parallel and Distributed Processing Symposium (IPDPS), 2017 IEEE International, IEEE, 2017, pp. 459–468.
- [6] P. Guo, H. Huang, Q. Chen, L. Wang, E.-J. Lee, P. Chen, A model-driven partitioning and auto-tuning integrated framework for sparse matrix-vector multiplication on gpus, in: Proceedings of the 2011 TeraGrid Conference: Extreme Digital Discovery, ACM, 2011, p. 2.
- [7] H. Huang, J. M. Dennis, L. Wang, P. Chen, A scalable parallel lsqr algorithm for solving large-scale linear system for tomographic problems: a case study in seismic tomography, *Procedia Computer Science* 18 (2013) 581–590.
- [8] H. Huang, L. Wang, E.-J. Lee, P. Chen, An mpi-cuda implementation and optimization for parallel sparse equations and least squares (lsqr), *Procedia Computer Science* 9 (2012) 76–85.
- [9] M. Zaharia, A. Konwinski, A. D. Joseph, R. H. Katz, I. Stoica, Improving mapreduce performance in heterogeneous environments., in: *Osdi*, Vol. 8, 2008, p. 7.
- [10] S. Lu, B. Rao, X. Wei, B. Tak, L. Wang, L. Wang, Log-based abnormal task detection and root cause analysis for spark, in: *Web Services (ICWS)*, 2017 IEEE International Conference on, IEEE, 2017, pp. 389–396.
- [11] S. Chauhan, P. Patel, F. C. Delicato, S. Chaudhary, A development framework for programming cyber-physical systems, in: *Proceedings of the 2nd International Workshop on Software Engineering for Smart Cyber-Physical Systems*, ACM, 2016, pp. 47–53.
- [12] Y. Nan, W. Li, W. Bao, F. C. Delicato, P. F. Pires, Y. Dou, A. Y. Zomaya, Adaptive energy-aware computation offloading for cloud of things systems, *IEEE Access* 5 (2017) 23947–23957.

- [13] G. Ananthanarayanan, S. Kandula, A. G. Greenberg, I. Stoica, Y. Lu, B. Saha, E. Harris, Reining in the outliers in map-reduce clusters using mantri., in: OSDI, Vol. 10, 2010, p. 24.
- [14] O. Ibidunmoye, F. Hernández-Rodriguez, E. Elmroth, Performance
675 anomaly detection and bottleneck identification, ACM Computing Surveys (CSUR) 48 (1) (2015) 4.
- [15] P. Garraghan, X. Ouyang, R. Yang, D. McKee, J. Xu, Straggler root-cause and impact analysis for massive-scale virtualized cloud datacenters, IEEE Transactions on Services Computing.
- [16] H. Jayathilaka, C. Krintz, R. Wolski, Performance monitoring and root
680 cause analysis for cloud-hosted web applications, in: Proceedings of the 26th International Conference on World Wide Web, International World Wide Web Conferences Steering Committee, 2017, pp. 469–478.
- [17] M. Y. Chen, E. Kiciman, E. Fratkin, A. Fox, E. Brewer, Pinpoint: Problem
685 determination in large, dynamic internet services, in: Dependable Systems and Networks, 2002. DSN 2002. Proceedings. International Conference on, IEEE, 2002, pp. 595–604.
- [18] X. Gu, H. Wang, Online anomaly prediction for robust cluster systems, in: Data Engineering, 2009. ICDE'09. IEEE 25th International Conference on,
690 IEEE, 2009, pp. 1000–1011.
- [19] E. W. Fulp, G. A. Fink, J. N. Haack, Predicting computer system failures using support vector machines., WASL 8 (2008) 5–5.
- [20] N. J. Yadwadkar, G. Ananthanarayanan, R. Katz, Wrangler: Predictable
695 and faster jobs using fewer resources, in: Proceedings of the ACM Symposium on Cloud Computing, ACM, 2014, pp. 1–14.
- [21] M. L. Massie, B. N. Chun, D. E. Culler, The ganglia distributed monitoring system: design, implementation, and experience, Parallel Computing 30 (7) (2004) 817–840.

- [22] A. Oliner, J. Stearley, What supercomputers say: A study of five system logs, in: DSN, IEEE, 2007.
- [23] S. Ryza, U. Laserson, S. Owen, J. Wills, Advanced Analytics with Spark: Patterns for Learning from Data at Scale, O'Reilly Media, 2015.
- [24] J. Tan, X. Pan, S. Kavulya, R. Gandhi, P. Narasimhan, Salsa: Analyzing logs as state machines., WASL 8 (2008) 6–6.
- [25] Q. Chen, D. Zhang, M. Guo, Q. Deng, S. Guo, Samr: A self-adaptive mapreduce scheduling algorithm in heterogeneous environment, in: Computer and Information Technology (CIT), 2010 IEEE 10th International Conference on, IEEE, 2010, pp. 2736–2743.
- [26] W. Xu, L. Huang, A. Fox, D. Patterson, M. I. Jordan, Detecting large-scale system problems by mining console logs, in: SOSP, ACM, 2009.
- [27] M. K. Aguilera, J. C. Mogul, J. L. Wiener, P. Reynolds, A. Muthitacharoen, Performance debugging for distributed systems of black boxes, ACM SIGOPS Operating Systems Review 37 (5) (2003) 74–89.
- [28] H. Zhou, Y. Li, H. Yang, J. Jia, W. Li, Bigroots: An effective approach for root-cause analysis of stragglers in big data system, arXiv preprint arXiv:1801.03314.
- [29] S. Diersen, E.-J. Lee, D. Spears, P. Chen, L. Wang, Classification of seismic windows using artificial neural networks, Procedia computer science 4 (2011) 1572–1581.
- [30] J. Shi, Y. Qiu, U. F. Minhas, L. Jiao, C. Wang, B. Reinwald, F. Özcan, Clash of the titans: Mapreduce vs. spark for large scale data analytics, Proceedings of the VLDB Endowment 8 (13) (2015) 2110–2121.
- [31] S. Huang, J. Huang, J. Dai, T. Xie, B. Huang, The hibenach benchmark suite: Characterization of the mapreduce-based data analysis, in: Data Engineering Workshops (ICDEW), 2010 IEEE 26th International Conference on, IEEE, 2010, pp. 41–51.